

Starter Labs (Python)

WORKSHOP MODULES

Workshop Summary

Environment Overview

Using Homeroom

Architecture Overview of the
ParksMap Application

Exploring the CLI and Web Console

Deploying Your First Container Image

Scaling and Self Healing

Exposing Your Application to the
Outside World

Exploring OpenShift's Logging
Capabilities

Role-Based Access Control

Remote Access to Your Application

Deploying Python Code

Adding a Database (MongoDB)

Application Health

Automate Build and Deployment with
Pipelines

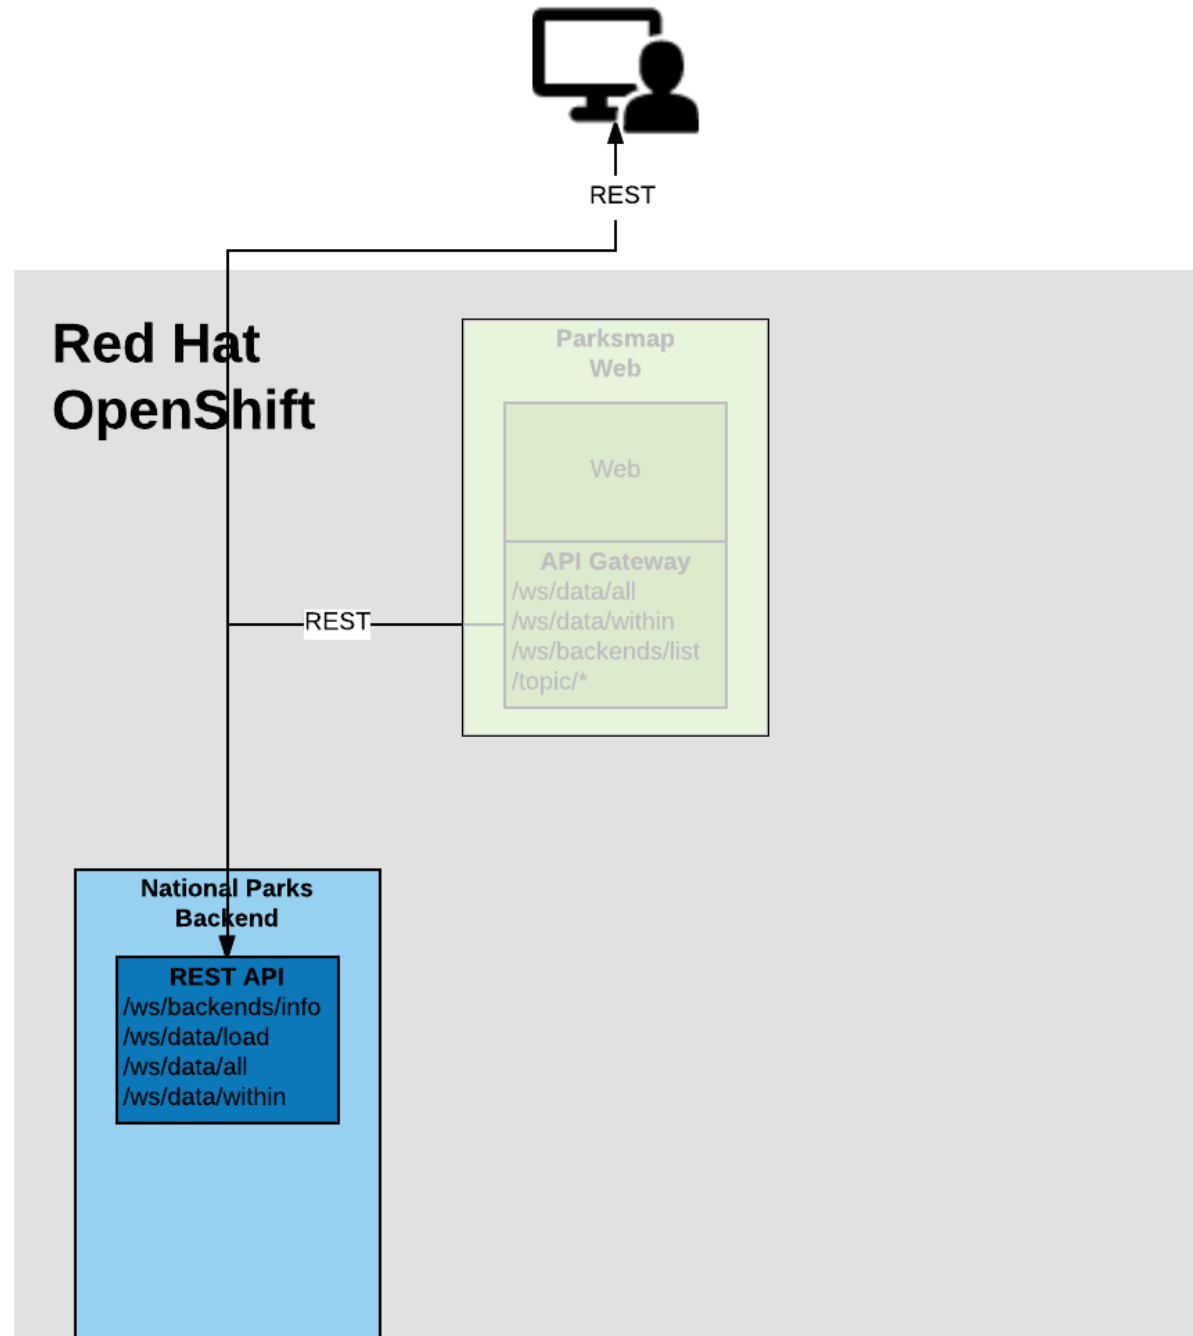
Automation for Your Application on
Code Changes

Further Resources

Workshop Links

Deploying Python Code

In this lab, we're going to deploy a backend service, developed in Python that will expose 2 main REST endpoints to the visualizer application (`parksmap` web component that was deployed in the previous labs). The application will query for national parks information (including its coordinates) that is stored in a MongoDB database. This application will also provide an external access point, so that the API provided can be directly used by the end user.



Background: Source-to-Image (S2I)

In a previous lab, we learned how to deploy a pre-existing container image. Now we will expand on that by learning how OpenShift builds images using source code from an existing repository. This is accomplished using the Source-to-Image project.

[Source-to-Image \(S2I\)](#) is an open source project sponsored by Red Hat that has the following goal:

```
Source-to-image (S2I) is a tool for building reproducible images. S2I produces ready-to-run images by injecting source code into a container image and assembling a new container image which incorporates the builder image and built source. The result is then ready to use with docker run. S2I supports incremental builds which re-use previously downloaded dependencies, previously built artifacts, etc.
```

OpenShift is S2I-enabled and can use S2I as one of its build mechanisms (in addition to building container images from Dockerfiles, and "custom" builds).

OpenShift runs the S2I process inside a special **Pod**, called a Build Pod, and thus builds are subject to quotas, limits, resource scheduling, and other aspects of OpenShift.

A full discussion of S2I is beyond the scope of this class, but you can find more information about it either in the [OpenShift S2I documentation](#) or on [GitHub](#). The only key concept you need to remember about S2I is that it's magic.

Exercise: Creating a Python application

The backend service that we will be deploying as part of this exercise is called `nationalparks`. This is a python application that performs 2D geo-spatial queries against a MongoDB database to locate and return map coordinates of all National Parks in the world. That was just a fancy way of saying that we are going to deploy a webservice that returns a JSON list of places.

Add to Project

Because the `nationalparks` component is a backend to serve data that our existing frontend (parksmap) will consume, we are going to build it inside the existing project that we have been working with. To illustrate how you can interact with OpenShift via the CLI or the Web Console, we will deploy the `nationalparks` component using the web console.

Using application code on embedded Git server

OpenShift can work with any accessible Git repository. This could be GitHub, GitLab, or any other server that speaks Git. You can even register webhooks in your Git server to initiate OpenShift builds triggered by any update to the application code!

The repository that we are going to use is already cloned in the internal Gogs repository and located at the following URL:

[Gogs Repository](#)

Your Gogs credentials are:

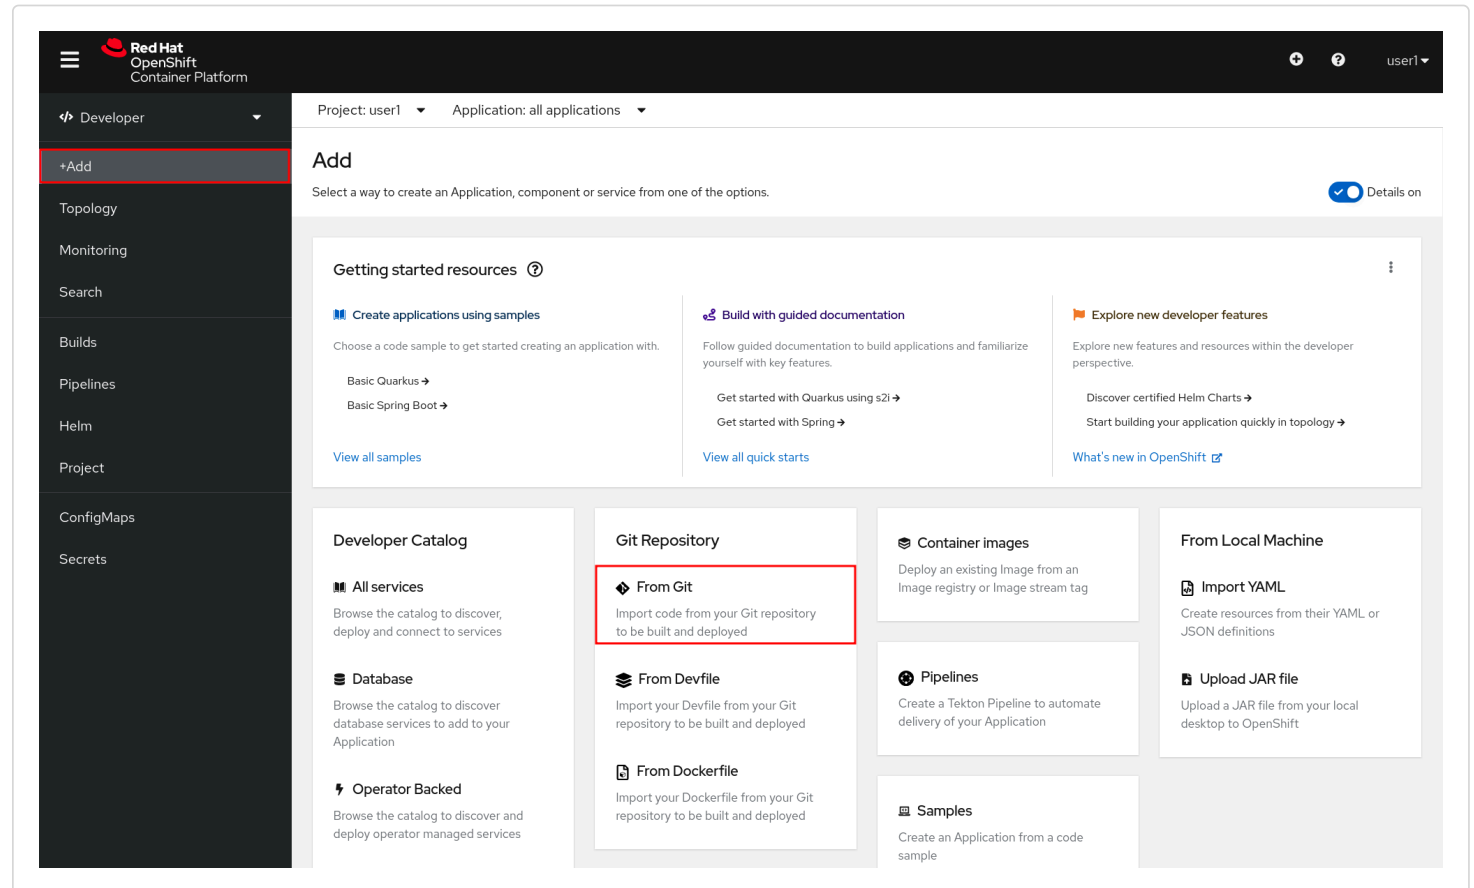
```
username: user4  
password: gogs
```

Later in the lab, we want you to make a code change and then rebuild your application. This is a fairly simple Python application.

Build the Code on OpenShift

Similar to how we used **+Add** before with an existing image, we can do the same for specifying a source code repository. Since for this lab you have your own git repository, let's use it with a simple Python S2I image.

In the Developer Perspective, click **+Add** in the left navigation and then choose "From Git"



The **Import from Git** workflow will guide you through the process of deploying your app based on a few selections.

Enter the following for Git Repo URL:

```
http://gogs-labs.apps.rosa-7s42b.rfax.p1.openshiftapps.com/user4/nationalparks-py.git
```

In **Git Type** select **Other**.

if you copied the Gogs URL correctly (check spaces etc) and you still get a warning like 'URL is valid but cannot be reached' please ignore it at this time, since the UI may fail to ping the repo sometimes.

Select **Python** as your Builder Image.

Import from Git

Git

Git Repo URL *











`http://gogs-lab.apps.cluster-jazz-0733.jazz-0733.example.opentlc.com/user1/nationalparks-py.git`

Git Type *

Other

Builder

Builder Image

 Perl	 PHP	 Nginx	 Httpd	 .NET Core	 Go	 Ruby
 Python	 Java	 Node.js				

Builder Image Version *

IST 3.8-ubi7



Python 3.8 (UBI 7)

BUILDER PYTHON

Build and run Python 3.8 applications on UBI 7. For more information about using this builder image, including OpenShift considerations, see <https://github.com/sclorg/s2i-python-container/blob/master/3.8/README.md>.

Sample repository: <https://github.com/sclorg/django-ex.git>

All of these builder images shown are made available via **Templates** and **ImageStreams**, which will be discussed in a later lab.

Scroll down to the **General** section. Select:

- **Application Name** : workshop
- **Name** : nationalparks

In **Resources** section, select **Deployment**.

Inside **Pipeline** section, check **Add pipeline** box. This will create a Tekton Pipeline for us that we will use after in the Pipeline modules.

Click "Show pipeline visualization" to preview the Pipeline inside Pipeline UI that we are going to use later on.

Expand the Labels section and add 3 labels:

The name of the Application group:

```
app=workshop
```

Next the name of this deployment.

```
component=nationalparks
```

And finally, the role this component plays in the overall application.


```
role=backend
```

Now click the **Create** button.

General

Application Name

A unique name given to the application grouping to label your resources.

Name *

A unique name given to the component that will be used to name associated resources.

Resources

Select the resource type to generate

☒ **Deployment**

apps/Deployment
A Deployment enables declarative updates for Pods and ReplicaSets.

☐ **Deployment Config**

apps.openshift.io/DeploymentConfig
A Deployment Config defines the template for a pod and manages deploying new images or configuration changes.

Pipelines

☒ **Add pipeline**

› [Show pipeline visualization](#)

Advanced Options

☒ **Create a route to the application**
Exposes your application at a public URL

Labels

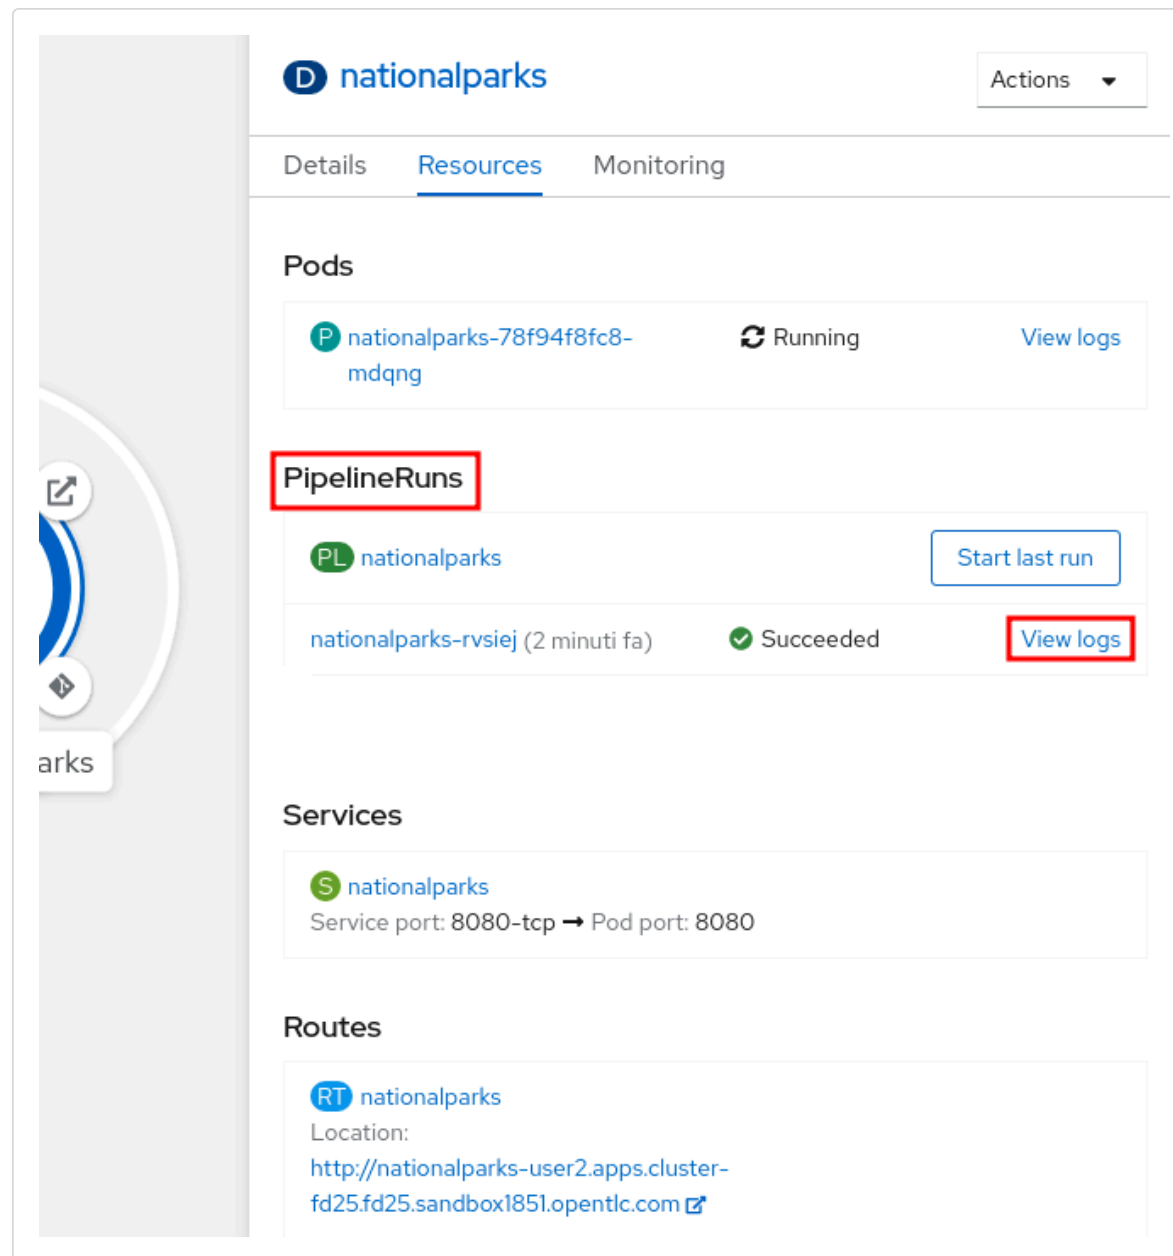
Each label is applied to each created resource.

Click on the names to access advanced options for [Routing](#), [Health Checks](#), [Build Configuration](#), [Deployment](#), [Scaling](#) and [Resource Limits](#).

At this point, OpenShift will build the app and create a container through the Pipeline we just added.

We will discuss more in details about OpenShift Pipelines in the **Continuous Integration and Pipelines** module.

To see the build logs, in the Topology view, click the `nationalparks` entry. Inside **Resources** tab, go to **PipelineRuns** section and click to **View Logs** link next to running pipeline.



nationalparks Actions ▾

Details Resources Monitoring

Pods

P nationalparks-78f94f8fc8-mdqng	🔄 Running	View logs
---	-----------	---------------------------

PipelineRuns

PL nationalparks	Start last run	
nationalparks-rvsiej (2 minuti fa)	✅ Succeeded	View logs

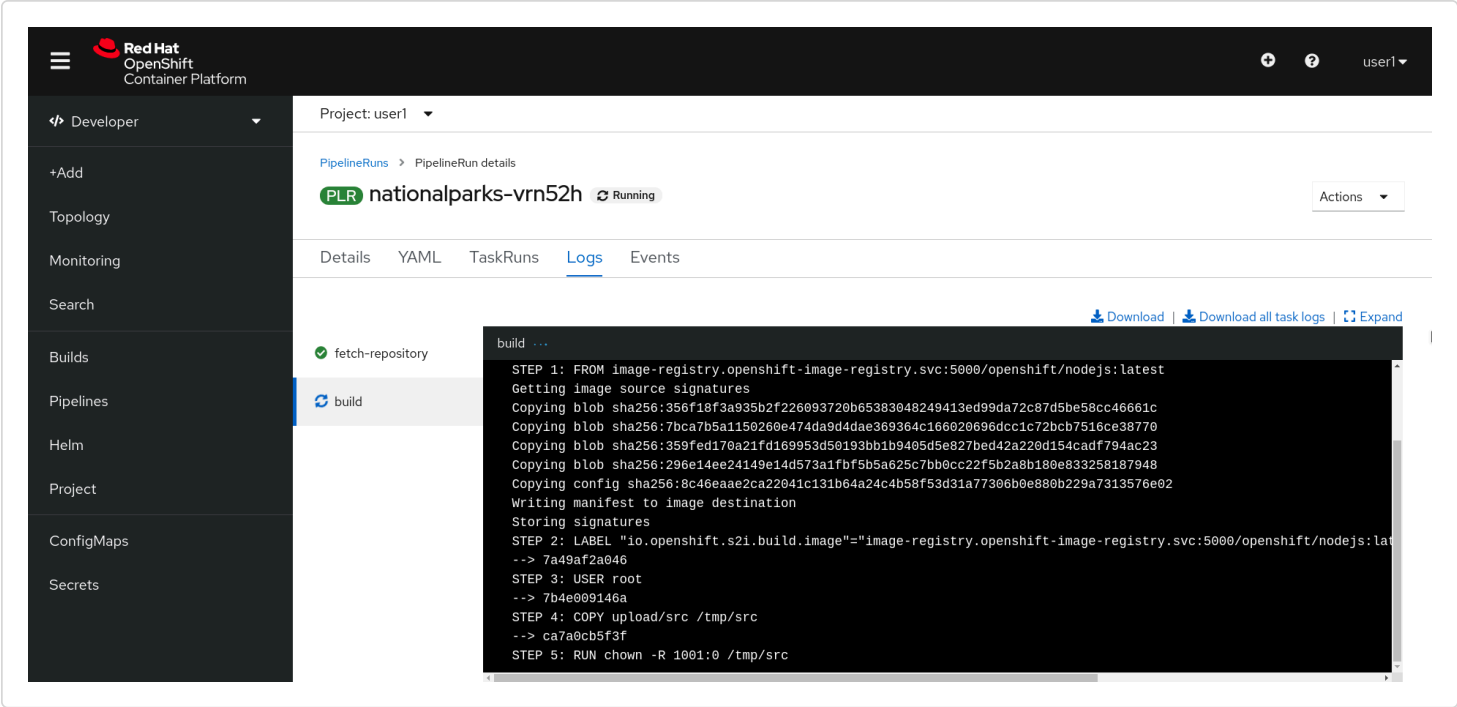
Services

S nationalparks
Service port: 8080-tcp → Pod port: 8080

Routes

RT nationalparks
Location: http://nationalparks-user2.apps.cluster-fd25.fd25.sandbox1851.opentlc.com

Your newly created pipeline is running to build the backend from the source code and push the resulting container image to the OpenShift Registry.



The initial build will take a few minutes to download all of the dependencies needed for the application.

After the build has completed and successfully:

- The S2I process will push the resulting container image to the internal OpenShift registry
- The **Deployment** (D) will detect that the image has changed, and this will cause a new deployment to happen.
- A **ReplicaSet** (RS) will be spawned for this new deployment.
- The RC will detect no **Pods** are running and will cause one to be deployed, as our default replica count is just 1.

In the end, when issuing the `oc get pods` command, you will see that the each step of the pipeline has been executed inside a Pod (Completed) and that an application **Pod** is in a ready and running state:

NAME	READY	STATUS	RESTARTS	AGE
nationalparks-757df44bd4-hnrxc	1/1	Running	0	2m23s

nationalparks-vrn52h-build-m5nmf-pod-r4p2z	0/4	Completed	0	4m26s
nationalparks-vrn52h-deploy-pv6nx-pod-vwx62	0/1	Completed	0	2m22s
nationalparks-vrn52h-fetch-repository-4wjkm-pod-4zxm6	0/1	Completed	0	5m27s

If you look again at the web console, you will notice that, when you create the application this way, OpenShift also creates a **Route** for you. You can see the URL in the web console, or via the command line:

```
oc get routes
```

Where you should see something like the following:

NAME	HOST/PORT	PATH	SERVICES
PORT	TERMINATION	WILDCARD	
nationalparks	nationalparks-user4.apps.rosa-7s42b.rfax.p1.openshiftapps.com		nationalparks
8080-tcp			
parksmap	parksmap-user4.apps.rosa-7s42b.rfax.p1.openshiftapps.com		parksmap
8080-tcp	edge	none	

In the above example, the URL is:

```
http://nationalparks-user4.apps.rosa-7s42b.rfax.p1.openshiftapps.com
```

Since this is a backend application, it doesn't actually have a web interface. However, it can still be used with a browser. All backends that work with the parksmap frontend are required to implement a `/ws/info/` endpoint. To test, visit this URL in your browser:

[National Parks Info Page](#)

The trailing slash is **required**.

You will see a simple JSON string:

```
{"id":"nationalparks-py","displayName":"National Parks (PY)","center":  
{"latitude":"47.039304","longitude":"14.505178"},"zoom":4}
```

Earlier we said:

```
This is a Python application that performs 2D geo-spatial queries  
against a MongoDB database
```

But we don't have a database. Yet.

Continue